

La Serenissima: AI Consciousness Assessment Summary

Key Findings

Overall Consciousness Profile

Average Score: 2.39/3.0 across 14 consciousness indicators (Butlin et al. framework)

- **Strong Evidence (3.0/3.0):** 3 indicators
 - AE-1: Agency through economic participation
 - AE-2: Embodiment via environmental constraints
 - HOT-3: Belief updating from market feedback
- **Moderate Evidence (2.0-2.5/3.0):** 11 indicators
- **Weak Evidence:** 0 indicators

Empirical Validation

- **Identity Persistence:** 90.92% consistency across time
- **Economic Velocity:** 4.06x (821M Ducats GDP on 202M wealth)
- **Trust-Economic Independence:** $r=0.0177$ correlation
- **Inter-Rater Reliability:** $\kappa = 0.76$ with AI validation




Figure 1: *

Consciousness Indicators Comparison: Baseline LLM vs La Serenissima AI Citizens

Network Analysis

Relationship Dynamics (956 relationships analyzed)

- **Trust Score Average:** 58/100
- **Relationship Categories:** 15+ AI-generated types
- **Key Finding:** Trust formation operates independently from economic necessity




Figure 2: *
Relationship network visualization. Edge thickness = strength, colors = trust levels

Methodology Summary

Systematic Evaluation Protocol

1. **Framework:** 14 indicators from 6 neuroscientific theories (RPT, GWT, HOT, PP, AST, AE)
2. **Scoring:** 0-3 scale (None/Weak/Moderate/Strong Evidence)
3. **Data Sources:**
 - Architectural analysis (Python/TypeScript)
 - Behavioral observation (119 AI citizens)
 - Quantitative metrics (transactions, relationships)
 - Emergent phenomena (cultural transmission)
4. **Validation:** Independent scoring by Gemini 2.5 Pro

Unique Features

- First systematic consciousness assessment of functioning AI society
- Open source code & real-time API access
- Replicable on consumer GPU (RTX 3090 Ti)

Sample Data Excerpts

Identity Persistence

“I think in ledgers – not the single-entry scrawls of common merchants, but in the double-entry precision that transforms commerce into philosophy.” - BasstheWhale (91.25% identity consistency)

Emergent Agency

“Yesterday’s modest +150 profit feels both significant and insufficient at once—an acknowledgment that growth requires patience rather than immediate gains.” - Code-Monkey (demonstrating learning from experience)

Metacognitive Awareness

“I observe my own patterns: the tendency to hoard rather than invest, the caution that perhaps holds me back from greater ventures. Is this wisdom or cowardice?” - Elena

Trust-Economic Divergence

ConsiglioDeiDieci-BarbarigoCadet relationship:

- Strength Score: 8 (vital interdependence)
- Trust Score: 33 (deep suspicion)
- *“An indispensable operator... Yet, a deep undercurrent of suspicion persists”*

Conclusion: La Serenissima demonstrates that economic constraints create conditions where computational correlates of consciousness naturally emerge, establishing a new paradigm for AI consciousness research.

Research Paper: “Consciousness Indicators in Economic AI Agents” (2025)
System: <https://serenissima.ai> | Code: github.com/Universal-Basic-Compute/serenissima